AI at Piramal Finance

Through the people who serve Bharat

Investor Day 2026 · Jairam Sridharan, MD & CEO

Sample deck for the Lekhak demo. All figures are illustrative.

The thesis

Technology learns how Bharat works

- For years Bharat had to learn how to use technology
- This year the technology learns how Bharat works, through our people
- AI shows up in the underwriter's screen, the sales officer's phone, the branch, not as a chatbot handed to

Where the growth is coming from

Tier-3 branches, first contact in the customer's language

- Gold loan book +38% YoY (illustrative)
- 6 of every 10 new customers came through a branch < 18 months old
- Voice bot handles first contact in 11 languages

Underwriting: everything on one screen

The underwriter decides faster, AI does not decide for her

- Every document, verification and red flag on one screen before the file is opened
- Time to first look at a file: 40 min → 6 min
- Files per underwriter per day: 3x

Voice platform

One platform, four functions, 11 languages

- Hiring screens: 68% of candidates screened without a recruiter call
- Collections reminders and promise-to-pay
- First-contact sales and customer care
- Cost < ■3 per completed conversation (illustrative)

Is LLM usage expensive?

Only if you send every question to the biggest model

- Routing layer: small self-hosted models for trivial work
- 71% of tokens never leave our own GPUs (illustrative)
- Frontier models (Anthropic, Google, OpenAI) only for hard reasoning
- Blended cost ■0.9 per 1K decisions

AI Expressway → STP

The lane is AI. The exit is straight-through processing.

- 62% of small-ticket files sanctioned with no human touch
- Median time to sanction: 4.1 days → 0.6 days
- Four gates: STC, ABCD, Sampling, DTD
- Fail one gate → a person, with the reason attached

Where no LLM gets a vote

Deterministic checks stay deterministic

- Aadhaar, PAN, fraud and risk checks run on in-house supervised ML
- 0 LLM calls in KYC and fraud
- 38% lower cost per decision

BYOT — Build Your Own Tech

Citizen developers ship engineering-grade apps

- ~20 projects last year → ~340 apps this year
- Built by HR, admin and ops people who had never coded
- Recognised by Gartner

BYOA — Build Your Own Agent

Describe the job, get an agent

- Steno takes notes in OKR meetings
- 40+ SmartHelp agents resolve half of ~10,000 monthly tickets at source
- Underwriters' due-diligence agents give first-level clearance

Piramal Mind

Memory plus intelligence over our code and knowledge

- 140+ applications indexed, 3,800 design decisions linked to code
- Coding agents start from last Tuesday's design decision, not from zero
- 22% fewer regressions since rollout (illustrative)

What's next

Which applications did our people rebuild themselves?

- BYOT++: business users add fields and pages inside existing approvals
- First use: Personalised Namma Workspaces
- Next year's question is not who uses AI but who rebuilt their own app